

# Multi-Domain Active Learning for Semi-Supervised Anomaly Detection

**Vincent Vercruyssen**

Lorenzo Perini

Wannes Meert

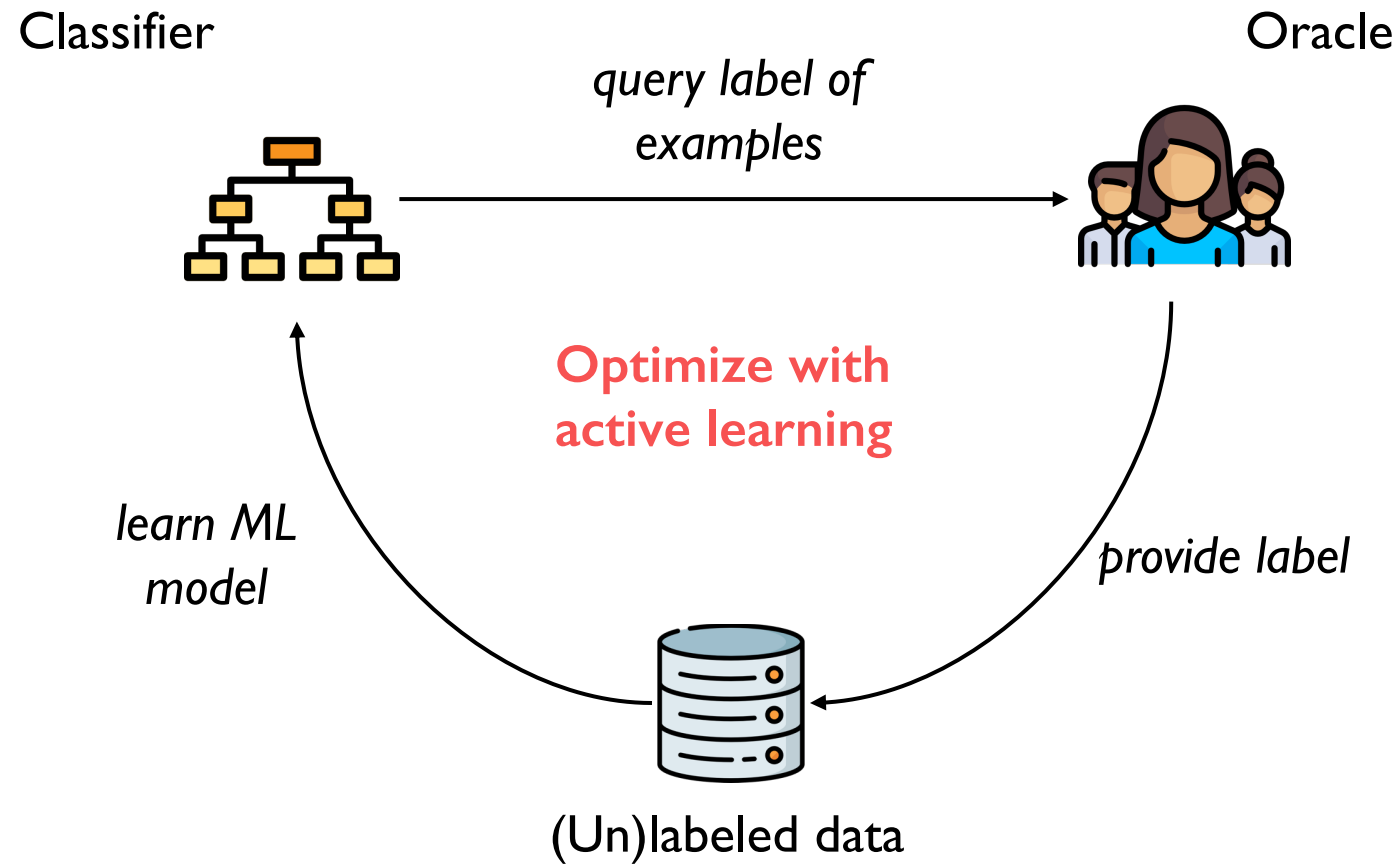
Jesse Davis

*DTAI Research Group, KU LEUVEN, Belgium*

[vincent.vercruyssen@kuleuven.be](mailto:vincent.vercruyssen@kuleuven.be)

@VercruyssenV

# Active learning





# Real-world problems pertain to multiple datasets

Water consumption is different in each store:

- On-site butchery or not
- Opening times
- Large or small store
- ...


...leading to different datasets!

Retail store




# Different datasets necessitate different models


Model 1



Model 2



Model 3






What if we only have 1 expert that has limited time to answer queries?

# I) Some datasets require more expert labels than others

---

**Labeling payoff** = marginal gain of acquiring another label for a given dataset



← LOW labeling payoff

→ HIGH labeling payoff

## 2) Active learning entails diminishing labeling payoff

The more labels given by the oracle for a dataset, the less the marginal gain of providing an additional label



# Multi-domain active learning

---

**GIVEN:** a multi-domain dataset  $M$  consisting of  $K$  datasets, a fixed query budget  $T$ , and an expert oracle  $O$

**DO:** learn a classifier for each of the  $K$  datasets with active learning


Pool-based active learning with a query batch size of one

Central to this problem is the **exploration** versus **exploitation** trade-off

↙  
= figuring out which  
datasets are most  
useful to label

↘  
= providing the most  
labels in datasets  
that need it most

# Exploration-exploitation with multi-armed bandits (MAB)







# An MAB algorithm to tackle multi-domain active learning

---

Labeling payoff  
distribution



Solving active learning via a multi-armed bandit approach is challenging:

1. How to equate groups of examples with the *arms* that the MAB can choose from?
2. How to select an *individual example* to query?
3. How do we quantify the *reward* for asking a query and getting the label?
4. How do we deal with the *diminishing returns* of labeling an additional example?

# I. An MAB algorithm chooses between *arms*

---

Each arm corresponds to selecting a *group of examples* from which to query one

1. Arm = one of the  $K$  *datasets* to query an example from

$$|arms| = K$$

✓ less exploration needed

✗ less fine-grained control

2. Arm = a *cluster in a dataset* to query an example from

$$|arms| = K \times C$$

✗ more exploration needed

✓ more fine-grained control

## 2. Randomly select the final query example from a group

---

The MAB algorithm chooses one arm at each iteration

Each arm corresponds to selecting a *group of examples*

1. Select an example in the group *randomly*
  - ✗ not necessarily query most informative examples
  - ✓ unbiased estimate of the reward distribution
2. Select an example in the group *heuristically*
  - ✓ select examples with highest estimated labeling payoff
  - ✗ distorts the reward distribution estimate

Reasoning in the paper!

### 3. Playing an arm results in an observed *reward*

---

The reward of each action taken by the MAB algorithm reflects the labeling payoff

1. Reward = *entropy reduction* in the predictions of the underlying classifier

$$r = \sum_{x \in D^k} [H_{f_+^k}(x) - H_{f^k}(x)]$$

2. Reward = *number of examples* for which the prediction of the classifier changes

$$r = 1 - \frac{Y_{f_+^k} \cdot Y_{f^k}}{\|Y_{f_+^k}\| \|Y_{f^k}\|}$$

## 4. Model diminishing labeling payoff with *rotting bandits*

The more labels given by the oracle for a dataset, the less the marginal gain of providing an additional label



ALBA uses the Sliding Window Average (SWA) rotting bandit algorithm

Levine, N., Crammer, K., and Mannor, S. (2017). Rotting bandits. In: *Advances in Neural Information Processing Systems* 30.



# Our *Active Learning Bandits* (ALBA) algorithm in full

---

**Input:** multi-domain dataset  $M$ , budget  $T$ , oracle  $O$ , number of clusters  $C$

**Output:** set of trained classifiers

1.  $A \leftarrow$  divide each dataset  $\in M$  into  $C$  clusters
2.  $F \leftarrow$  Train an initial classifier for each dataset  $\in M$
3. WHILE  $t < T$ :
  4.  $k \leftarrow$  SWA picks a cluster  $\in A$  based on the reward distribution
  5.  $x \leftarrow$  randomly select example in cluster  $k$
  6.  $l \leftarrow$  query  $x$  and receive label from  $O$
  7. update the classifier for the corresponding dataset
  8. compute reward and update cluster  $k$ 's reward distribution
  9.  $t = t + 1$

# Differences between ALBA and active learning

---

## Classic active learning

1. Estimates labeling payoff of every *single* example
2. Estimates labeling payoff of an example *before* it is queried
3. Works for a single dataset with a single classifier

## Active learning Bandits (ALBA)

1. Estimates labeling payoff of *groups* of examples
2. Directly observes true payoff *after* an example is queried
3. Works for multiple datasets, each with its own classifier

# Experimental evaluation

---

Our paper addresses three research questions:

- **Q1:** does ALBA outperform the existing active learning baselines for multi-domain active learning
- **Q2:** how does the division in groups of examples impact ALBA's performance
- **Q3:** how do the choice of reward function and query selection strategy impact ALBA's performance



# Benchmark based on retail water consumption data

---



Retail store

Several years of [water consumption data](#) for 7 retail stores

Task = classifying periods of abnormal water consumption


Final benchmark of 54 multi-domain datasets



# Benchmark based on retail water consumption data

---

From the 7 datasets, we constructed a full benchmark of 54 multi-domain datasets





# Baselines and evaluation

---

## 7 baselines

Combine all datasets into 1 dataset and learn a single classifier:

- **C-RAND**: acquire new labels randomly
- **C-UC**: acquire new labels heuristically

Treat each dataset independently and learn a separate classifier for each dataset:

- **I-U**: acquire no labels
- **I-RAND**: acquire new labels randomly
- **I-UC**: acquire new labels heuristically
- **I-R-RAND**: I-RAND + max queries / dataset
- **I-R-UC**: I-UC + max queries / dataset

## Evaluation


*Area under the ROC curve* measures the base performance of each anomaly classifier (we use the SSDO detector)

AUROC averaged over the K classifiers

*The learning curve* plots the AUROC as a function of the number of queried & labeled examples

*Area under the learning curve (AULC)* measures the impact of the active learning strategy (higher = better)

# Q1: ALBA outperforms the active learning baselines (1/2)



Friedman test: not all methods perform similarly

Bonferroni-Dunn test: ALBA is significantly better @ 100 query rounds

# QI: ALBA outperforms the active learning baselines (2/2)

Method	Nr. of times ALBA:			Ranking
	wins	draws	loses	Avg. $\pm$ SD
ALBA	-	-	-	<b>1.315 <math>\pm</math> 0.894</b>
I-RAND	<b>48</b>	2	4	2.639 $\pm$ 0.573
I-R-RAND	<b>48</b>	2	4	2.639 $\pm$ 0.573
I-U	<b>48</b>	2	4	4.333 $\pm$ 1.656
I-UC	<b>53</b>	0	1	5.157 $\pm$ 0.551
I-R-UC	<b>53</b>	0	1	5.231 $\pm$ 0.497
C-RAND	<b>54</b>	0	0	6.741 $\pm$ 0.865
C-UC	<b>54</b>	0	0	7.944 $\pm$ 0.404

(a) Results @ 100 query rounds

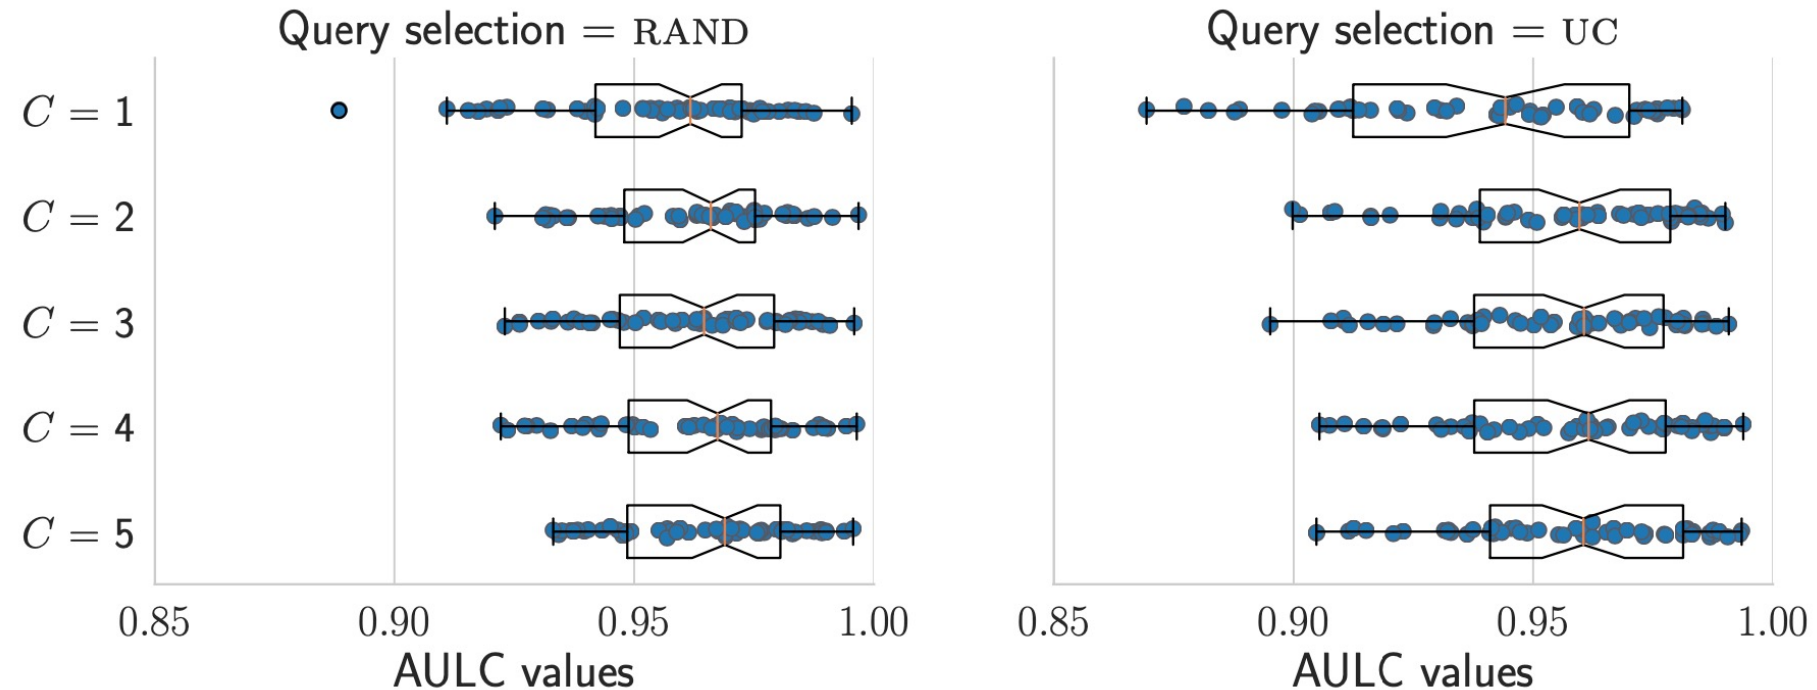
Method	Nr. of times ALBA:			Ranking
	wins	draws	loses	Avg. $\pm$ SD
ALBA	-	-	-	<b>1.306 <math>\pm</math> 0.710</b>
I-RAND	<b>45</b>	2	7	2.417 $\pm$ 0.507
I-R-RAND	<b>46</b>	1	7	2.417 $\pm$ 0.507
I-R-UC	<b>54</b>	0	0	4.537 $\pm$ 0.686
I-UC	<b>53</b>	1	0	4.546 $\pm$ 0.512
I-U	<b>54</b>	0	0	6.370 $\pm$ 0.818
C-RAND	<b>54</b>	0	0	6.491 $\pm$ 0.717
C-UC	<b>54</b>	0	0	7.917 $\pm$ 0.382

(b) Results @ 500 query rounds

Friedman test: not all methods perform similarly

Bonferroni-Dunn test: ALBA is significantly better @ 100 query rounds

## Q2: More clusters improves ALBA's performance



Correlation: higher  $C$  = higher ALBA performance

## Q3: cosine reward function + random example selection

---

Reward function	Query sel. strategy	Ranking Avg. $\pm$ SD
cosine	rand	<b>1.806 <math>\pm</math> 0.813</b>
cosine	uc	2.944 $\pm$ 0.926
entropy	rand	2.241 $\pm$ 0.843
entropy	uc	3.009 $\pm$ 0.825

Reward function: cosine  $>$  entropy

Example selection strategy: random  $>$  heuristic (uncertainty sampling)



# Conclusions

---

1. Multi-armed bandit strategies are an effective tool for multi-domain active learning
2. Our contribution, the ALBA algorithm, outperforms existing AL algorithms
3. Constructing a large benchmark helped us gain insight in the performance

Feel free to ask any question!

Vincent Vercruyssen, *Lorenzo Perini, Wannes Meert, Jesse Davis*

<https://github.com/Vincent-Vercruyssen/ALBA-paper>